




From Shallow to Deep Neural Networks

[Building More Complex Models]

José Oramas

Previous Session: Artificial Neurons

An artificial counterpart to real neurons



$$\sum_{i=1}^d w_i x_i + b$$
$$\sum_{i=0}^d w_i x_i, \quad x_0 := 1$$

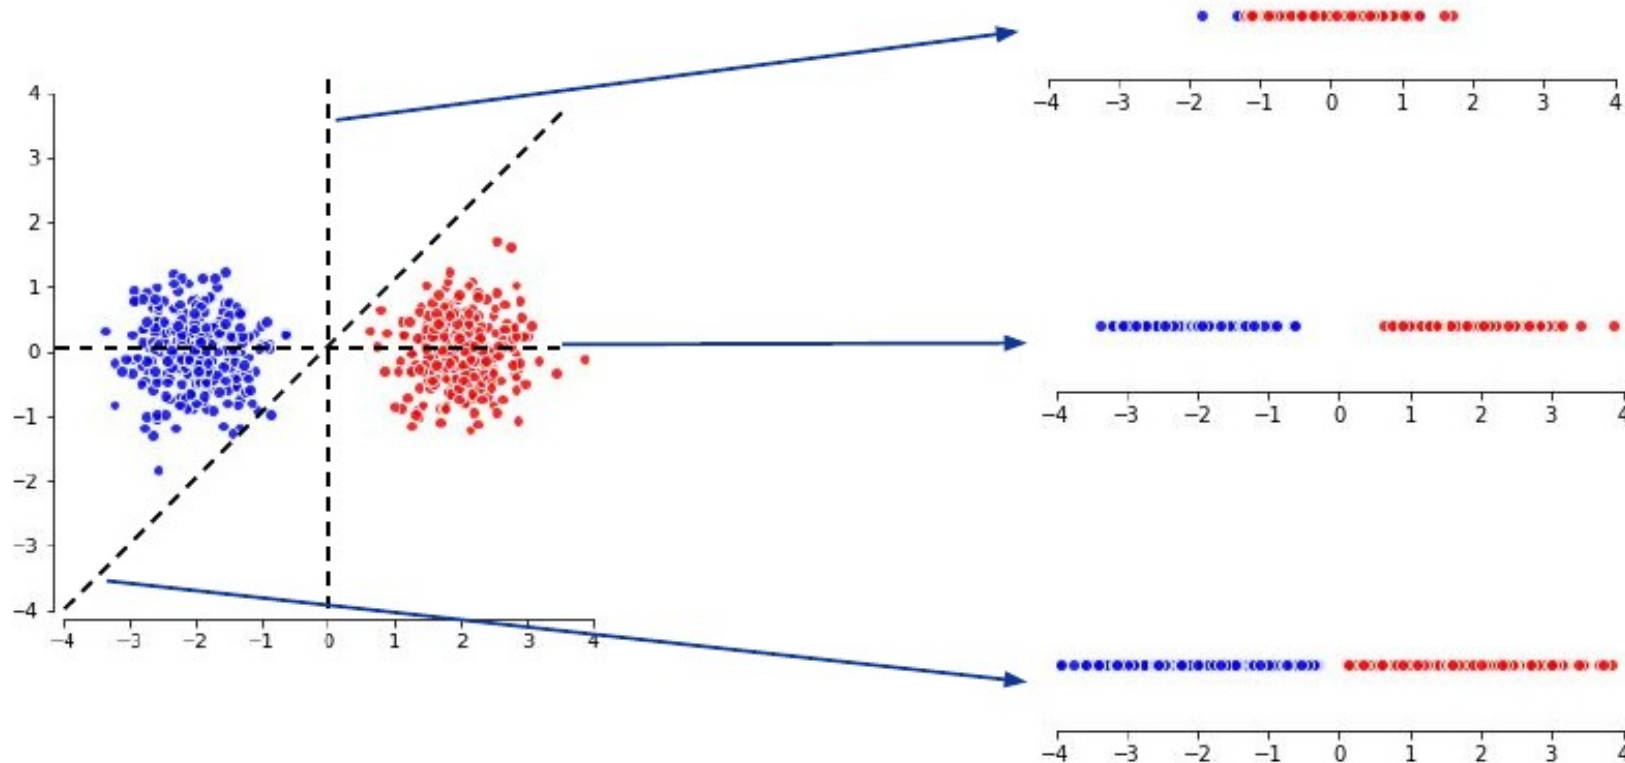
Characteristics:

- Basic computation
- Has inhibition/excitation connections
- Building block
- Time-independent state
- Outputs real values

Previous Session: Artificial Neurons

What they do?

- Define a linear (afine) projection of the data



Oh yes I remember, but ...
How do we use that?




A Shallow Neural Network

[with few layers]

From Neurons to Layers

A Common Composition


- Add several neurons working on “parallel”.



From Neurons to Layers

A Common Composition


- Add several neurons working on “parallel”.



From Neurons to Layers

A Common Composition


- Add several neurons working on “parallel”.



From Neurons to Layers

A Common Composition


- Add several neurons working on “parallel”.



From Neurons to Layers

A Common Composition


- Add several neurons working on “parallel”.



From Neurons to Layers

A Common Composition

- Add several neurons working on “parallel”.




$$h(x, w, b) = \langle w, x \rangle + b$$

$$f_{linear}(x, W, b) = Wx + b$$

From Neurons to Layers

A Common Composition

- Add several neurons working on “parallel”.



$$h(x, w, b) = \langle w, x \rangle + b$$

$$f_{linear}(x, W, b) = Wx + b$$


Why?

- Highly optimizable
 - Algorithmically (via smart matrix multiplication)
 - Hardware-wise (via GPUs, TPUs)
- Enable powerful compositions

From Layers to Neurons

Enabling Powerful Composition

- Add several neurons working on “parallel”.



Idea:

- Every neuron/layer → simple operation
- Using simple operations to build more complex ones
- Obtain a new quality out of the composition


A Single-Layer Neural Network

[with few layers]

A Single Layer Neural Network

Given:


- Let's assume we have the following simple model



A Single Layer Neural Network

Given:

- Let's assume we have the following simple model



Ok I see where this goes, but ...
**How do we train one
of those networks?**




Ok I see where this goes, but ...
**How do we train a
machine learning
model in general?**



Training a Model

Given:

- *Classification Task* with k classes.
- Training Data: inputs (x_i) and labels (\hat{y}_i)



A Simple Neural Network

How do we train such a model? → Learn the weights

- Let's see how we did it earlier

Algorithm: Perceptron Learning Algorithm

$P \leftarrow \text{inputs with label } 1;$

$N \leftarrow \text{inputs with label } 0;$

Initialize \mathbf{w} randomly;

while !convergence **do**

 Pick random $\mathbf{x} \in P \cup N$;

if $\mathbf{x} \in P$ and $\mathbf{w} \cdot \mathbf{x} < 0$ **then**

$\mathbf{w} = \mathbf{w} + \mathbf{x}$;

end

if $\mathbf{x} \in N$ and $\mathbf{w} \cdot \mathbf{x} \geq 0$ **then**

$\mathbf{w} = \mathbf{w} - \mathbf{x}$;

end

end

//the algorithm converges when all the
inputs are classified correctly


Requirements:

- Examples (with labels)
- A way to evaluate the “goodness” of the model
(measure performance)
- Stopping criteria

A Single-Layer Neural Network

Given:


- Let's assume we have the following simple model



A Single-Layer Neural Network

Given:


- Let's assume we have the following simple model



A Single-Layer Neural Network

Given:

- Let's assume we have the following simple model



$$h(x, w, b) = \langle w, x \rangle + b$$


$$f_{linear}(x, W, b) = Wx + b$$

The very same equations of
layers of artificial perceptrons

A Single-Layer Neural Network

Given:


- Let's assume we have the following simple model



A Single-Layer Neural Network


Given:

- Let's assume we have the following simple model



A Single-Layer Neural Network


Activation Function – Sigmoid



$$f_{\sigma}(x) = \frac{1}{1+e^{-x}} = \frac{e^x}{e^x+1}$$

A Single-Layer Neural Network

Activation Function – Sigmoid



Characteristics:

- Introduces non-linear behavior
- Scaled output [0-1]
- Simple derivatives
- Saturates
 - Vanishing derivatives

Note:


- Often called “non-linearities”
- Applied point-wise

$$f_{\sigma}(x) = \frac{1}{1+e^{-x}} = \frac{e^x}{e^x+1}$$

A Single-Layer Neural Network

Given:


- Let's assume we have the following simple model



A Single-Layer Neural Network

Given:


- Let's assume we have the following simple model



A Single-Layer Neural Network



Given:

- Let's assume we have the following simple model



A Single-Layer Neural Network

Loss Function – Cross Entropy



$$\ell_{CE}(p, t) = -[t \log(p) + (1 - t) \log(1 - p)]$$


Characteristics:

- Negation of logarithm of probability of correct prediction
- Composable with sigmoid
- Numerically unstable

A Single-Layer Neural Network

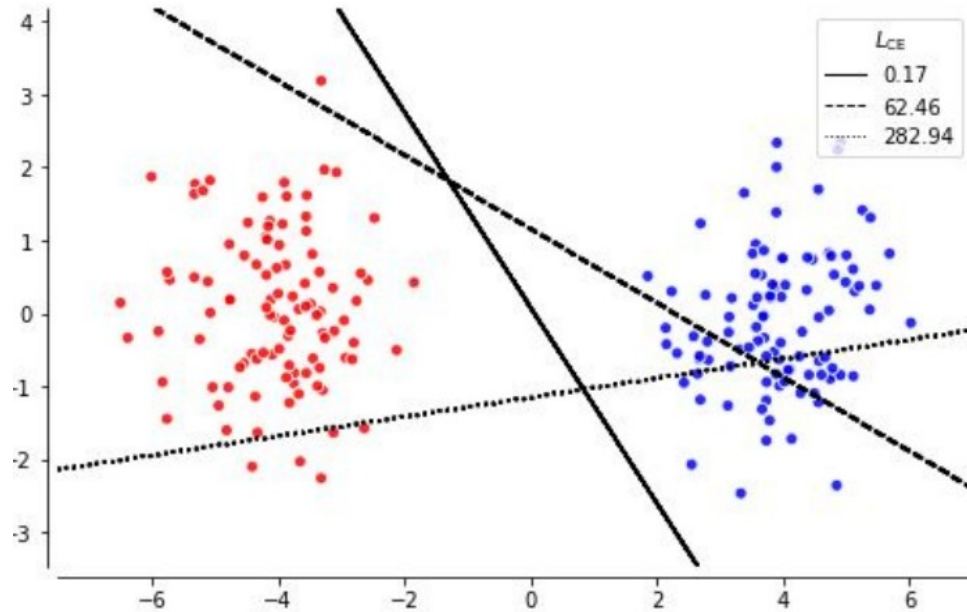
Given:

- Let's assume we have the following simple model



A Single-Layer Neural Network

Loss Function – Cross Entropy



Characteristics:

- **Additive w.r.t. samples**
(highly desirable)
- Negation of logarithm of probability of correct prediction
(on the entire dataset)
- Numerically unstable

$$L_{CE}(p, t) = - \sum_{i=1}^n [t^{(i)} \log(p^{(i)}) + (1 - t^{(i)}) \log(1 - p^{(i)})]$$

Beyond Binary Classification

[with few layers + multiple classes are possible]

A Single-Layer Neural Network

Beyond Binary Classification – Softmax

$$f_{sm}(x) = \frac{e^{x_i}}{\sum_{j=1}^k e^{x_j}}$$

A Single-Layer Neural Network

Beyond Binary Classification – Softmax

$$f_{sm}(x) = \frac{e^x}{\sum_{j=1}^k e^{x_j}}$$

Considering,

$$\begin{aligned} f_{sm}([x, 0]) &= \left[\frac{e^x}{e^x + e^0}, \frac{e^0}{e^x + e^0} \right] \\ &= [f_{\sigma}(x), 1 - f_{\sigma}(x)] \end{aligned}$$

A Single-Layer Neural Network

Beyond Binary Classification – Softmax

$$f_{sm}(x) = \frac{e^x}{\sum_{j=1}^k e^{x_j}}$$

Considering,

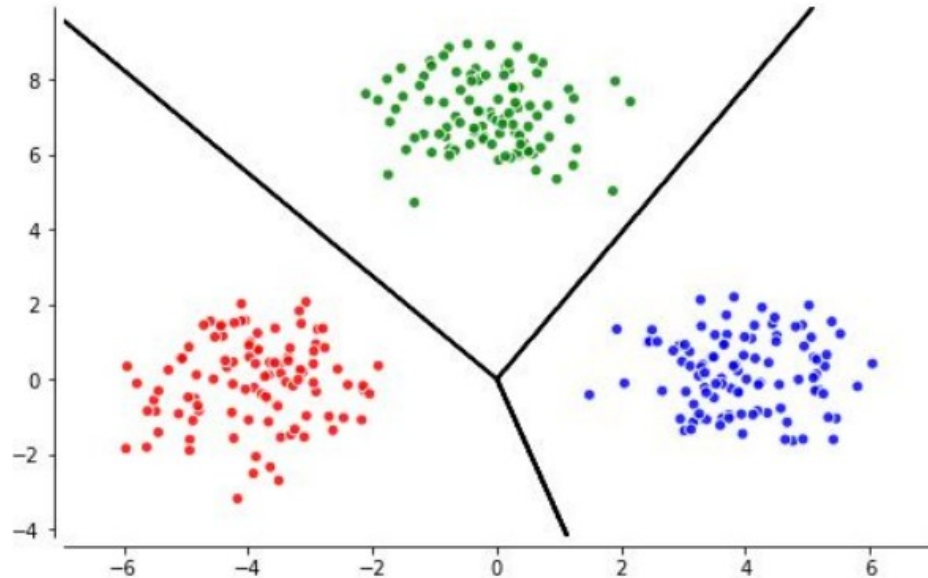
$$\begin{aligned} f_{sm}([x, 0]) &= \left[\frac{e^x}{e^x + e^0}, \frac{e^0}{e^x + e^0} \right] \\ &= [f_{\sigma}(x), 1 - f_{\sigma}(x)] \end{aligned}$$

Characteristics:

- Generalization of the sigmoid
- Does not work properly with sparse outputs
- Does not scale properly w.r.t. the number of classes (k)

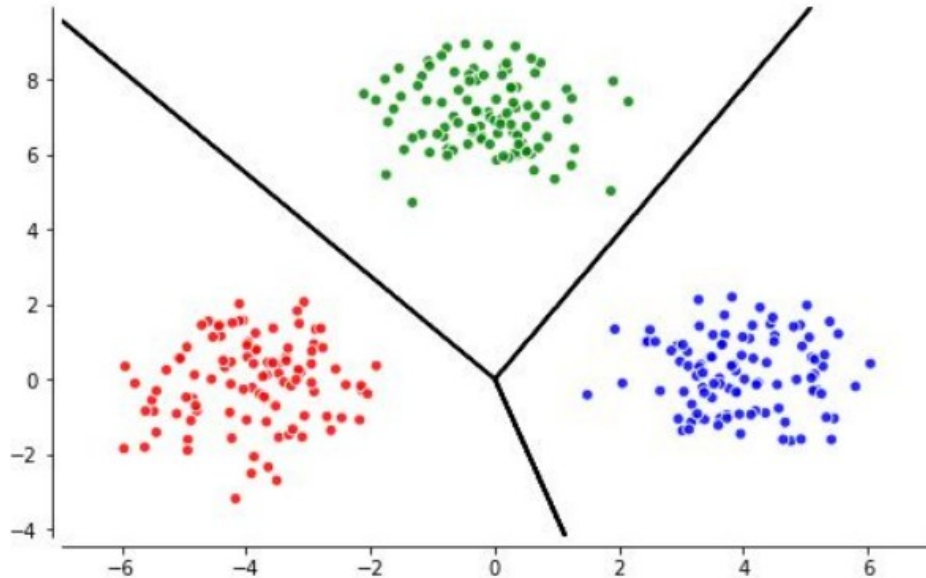
A Single-Layer Neural Network

Beyond Binary Classification – Softmax + Cross Entropy



A Single-Layer Neural Network


Beyond Binary Classification – Softmax + Cross Entropy



$$l_{CE}(f_{sm}(\mathbf{x}), \mathbf{t}) = - \sum_{j=1}^k t_j \log[f_{sm}(\mathbf{x}_j)]$$

A Single-Layer Neural Network

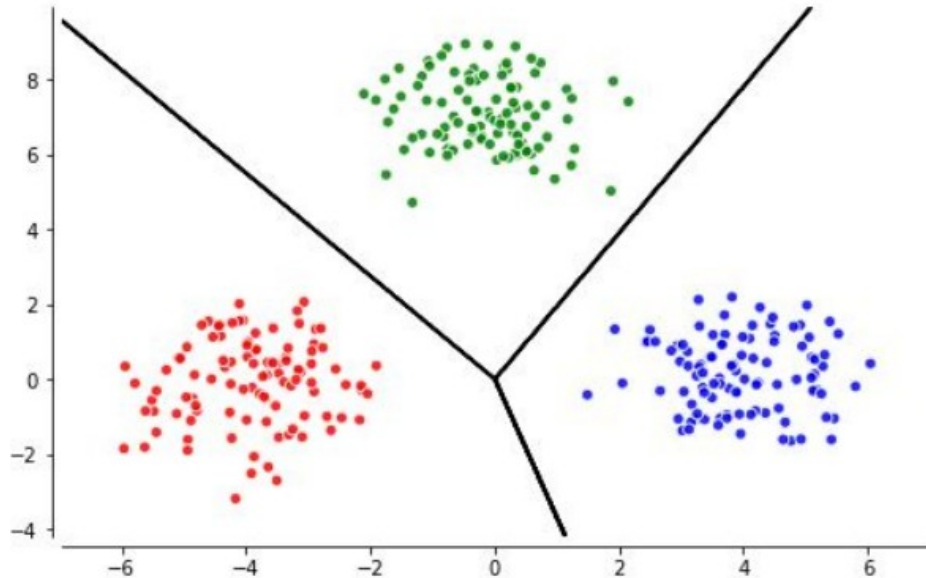
Beyond Binary Classification – Softmax + Cross Entropy



$$l_{CE}(f_{sm}(\mathbf{x}), \mathbf{t}) = - \sum_{j=1}^k t_j \log[f_{sm}(\mathbf{x}_j)] = - \sum_{j=1}^k t_j [\mathbf{x}_j - \log \sum_{l=1}^k e^{\mathbf{x}_l}]$$

A Single-Layer Neural Network

Beyond Binary Classification – Softmax + Cross Entropy



Characteristics:


- Generalization of the sigmoid
- Becomes numerically stable
- Simple, yet powerful
(~92% in handwritten digit recognition)

$$l_{CE}(f_{sm}(x), t) = - \sum_{j=1}^k t_j \log[f_{sm}(x_j)] = - \sum_{j=1}^k t_j [x_j - \log \sum_{l=1}^k e^{x_l}]$$

A Single-Layer Neural Network

Beyond Binary Classification – Softmax + Cross Entropy

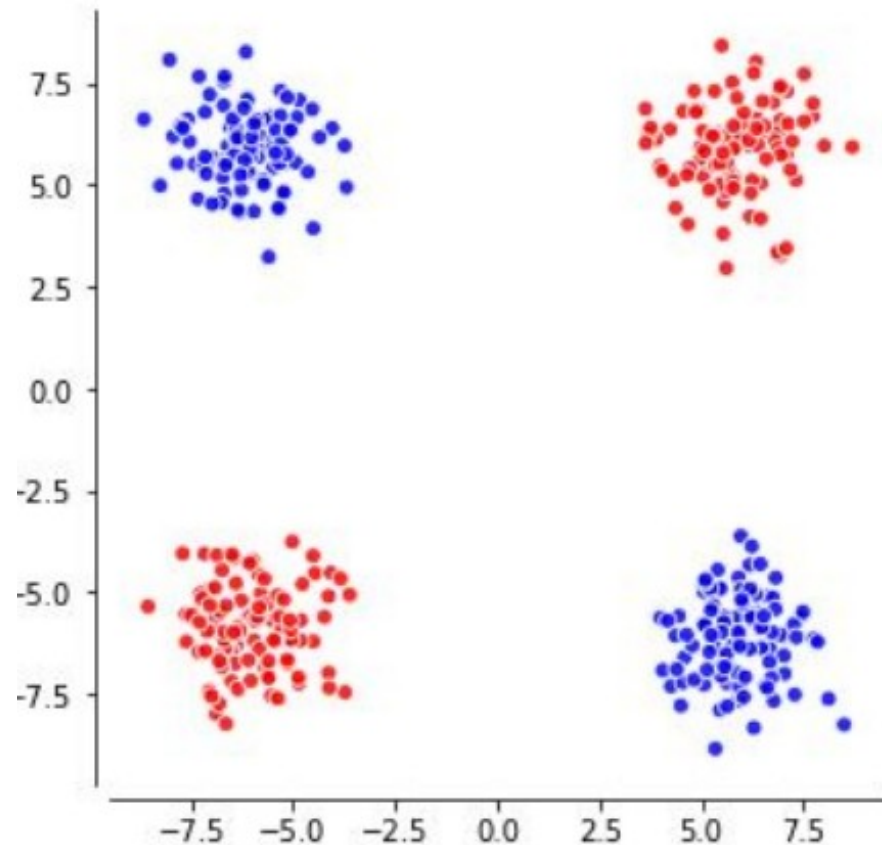
- Simple, yet powerful (~92% in handwritten digit recognition)



MNIST dataset, [Le Cunn et al., 1998a]

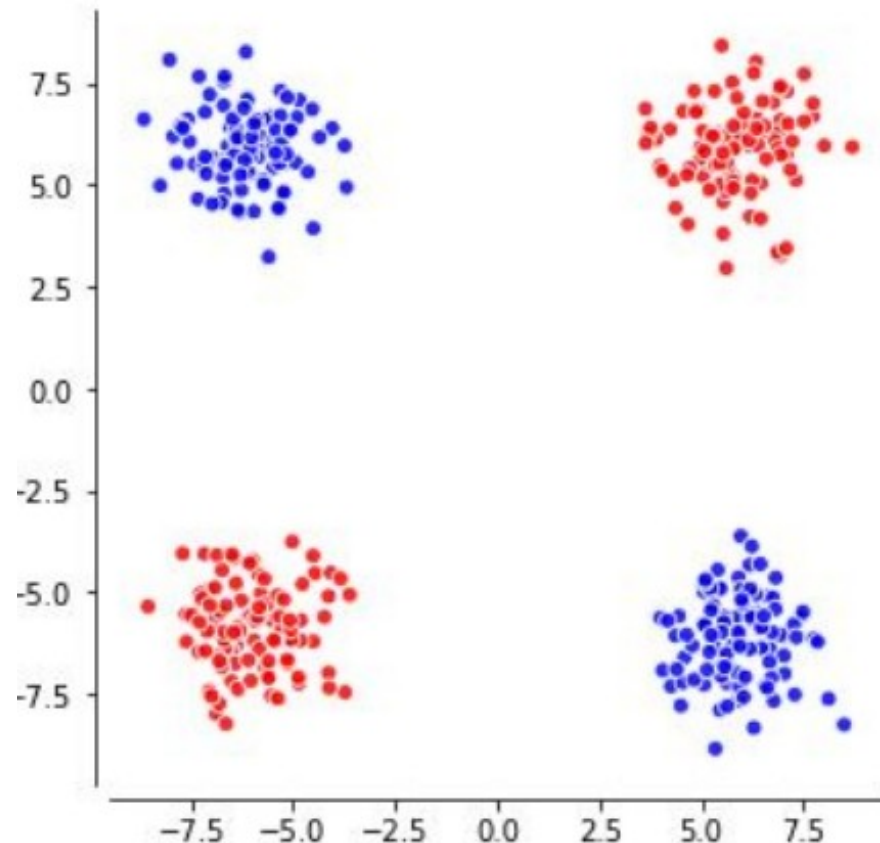
A Single-Layer Neural Network

What if we encounter the following problem?



A Single-Layer Neural Network

What if we encounter the following problem?



Exclusive Disjunction (XOR)

$$p \oplus q = (p \vee q) \wedge \neg(p \wedge q)$$

A Two-Layer Neural Network

[with few layers]

Break


See you in few minutes

A Two-Layer Neural Network

[with few layers]


A Two-Layer Neural Network

Extending the previous schematic



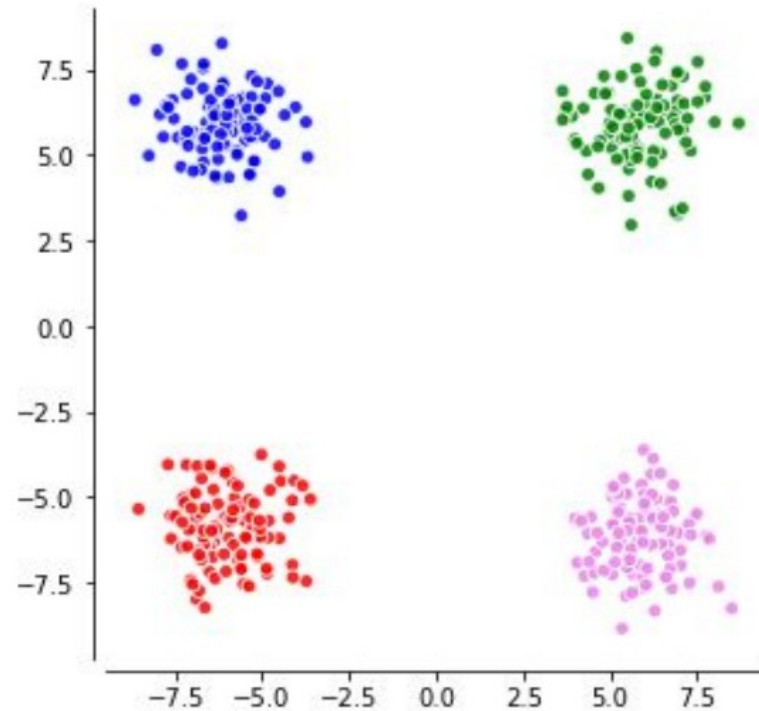
A Two-Layer Neural Network

Going back to the original problem




A Two-Layer Neural Network

Going back to the original problem



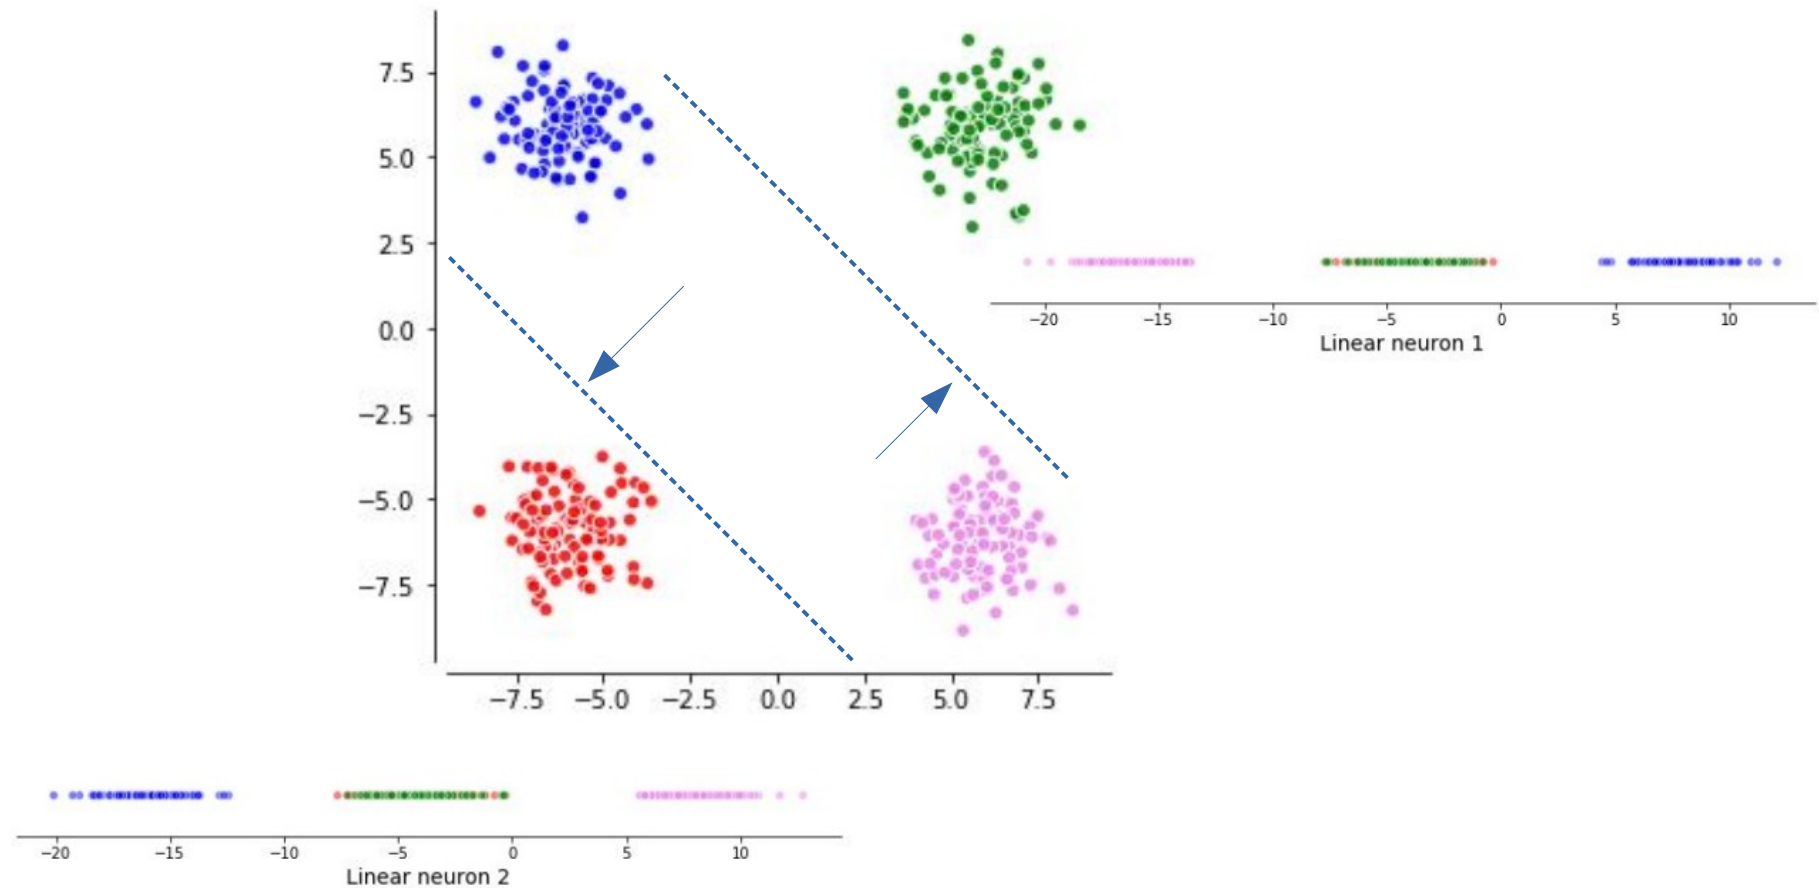
A Two-Layer Neural Network

Let's assume we have the following hyperplanes



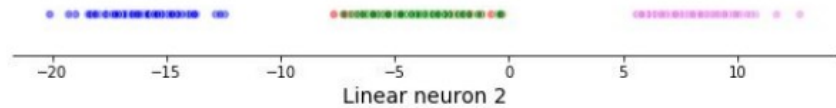
A Two-Layer Neural Network

Projecting our samples on the planes




A Two-Layer Neural Network

Going back to the original problem

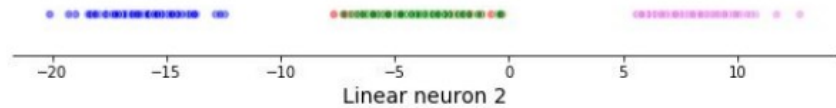


- Squashing our samples to the range $[0,1]$

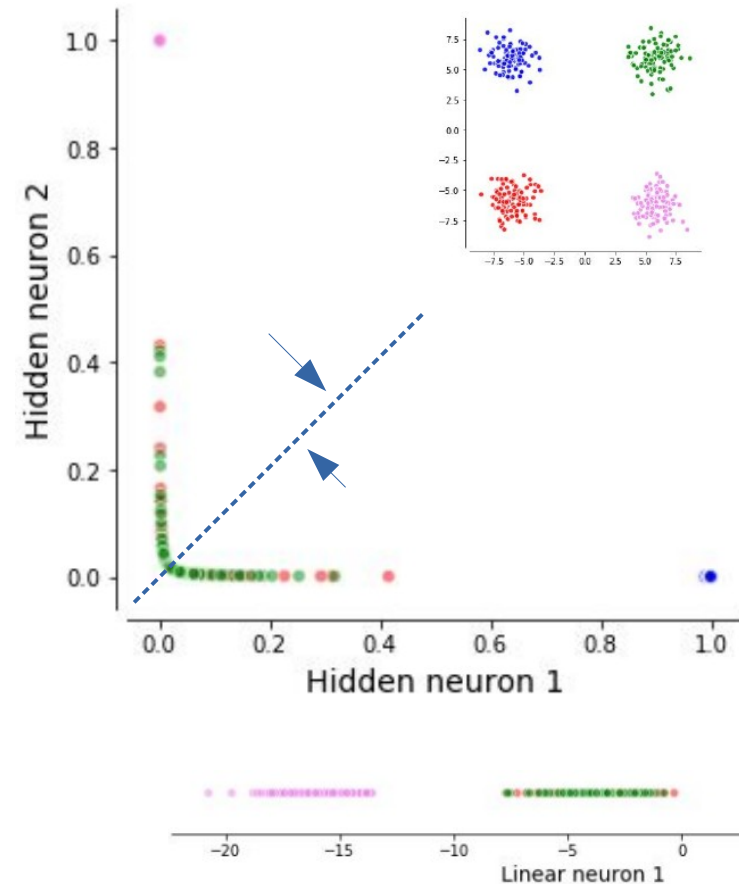


A Two-Layer Neural Network

Going back to the original problem

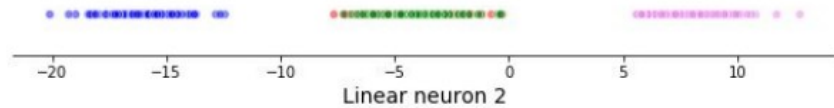


- Squashing our samples to the range $[0,1]$

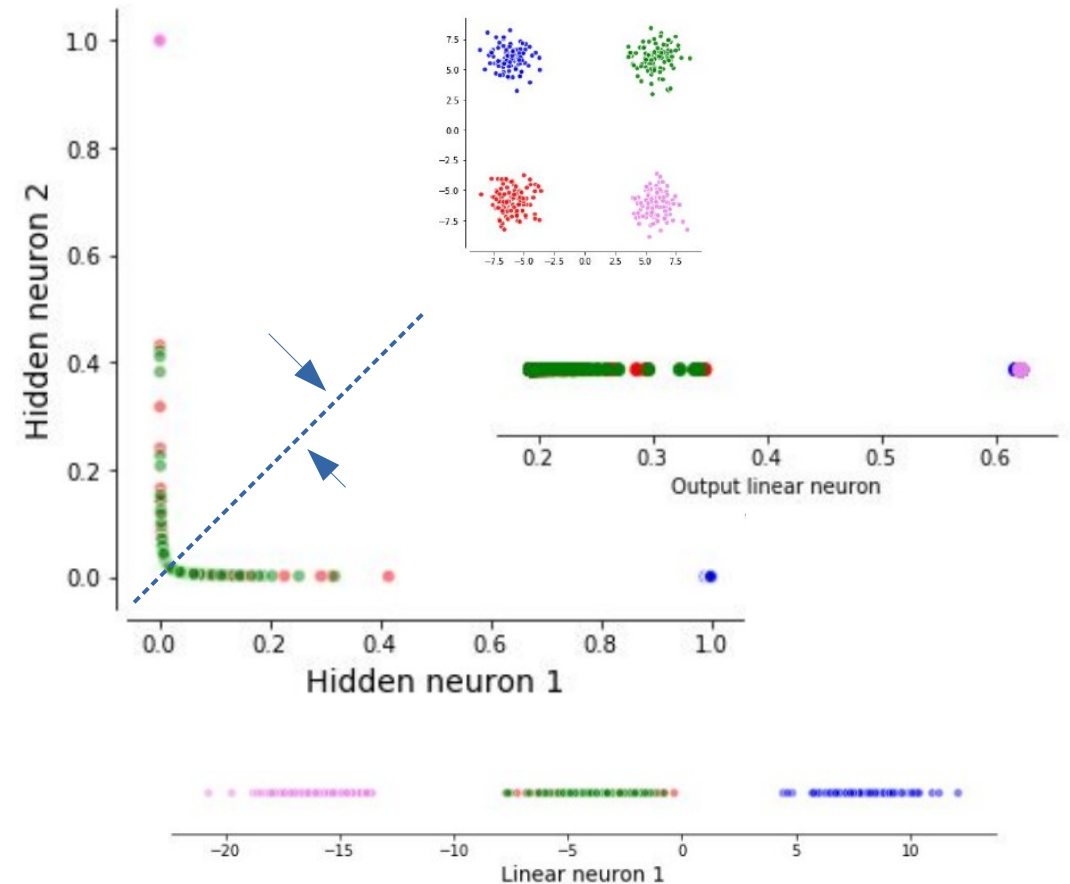


A Two-Layer Neural Network

Going back to the original problem

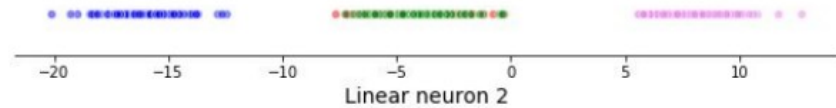


- Squashing our samples to the range $[0,1]$
- The hidden-layer provides a non-linear input space.

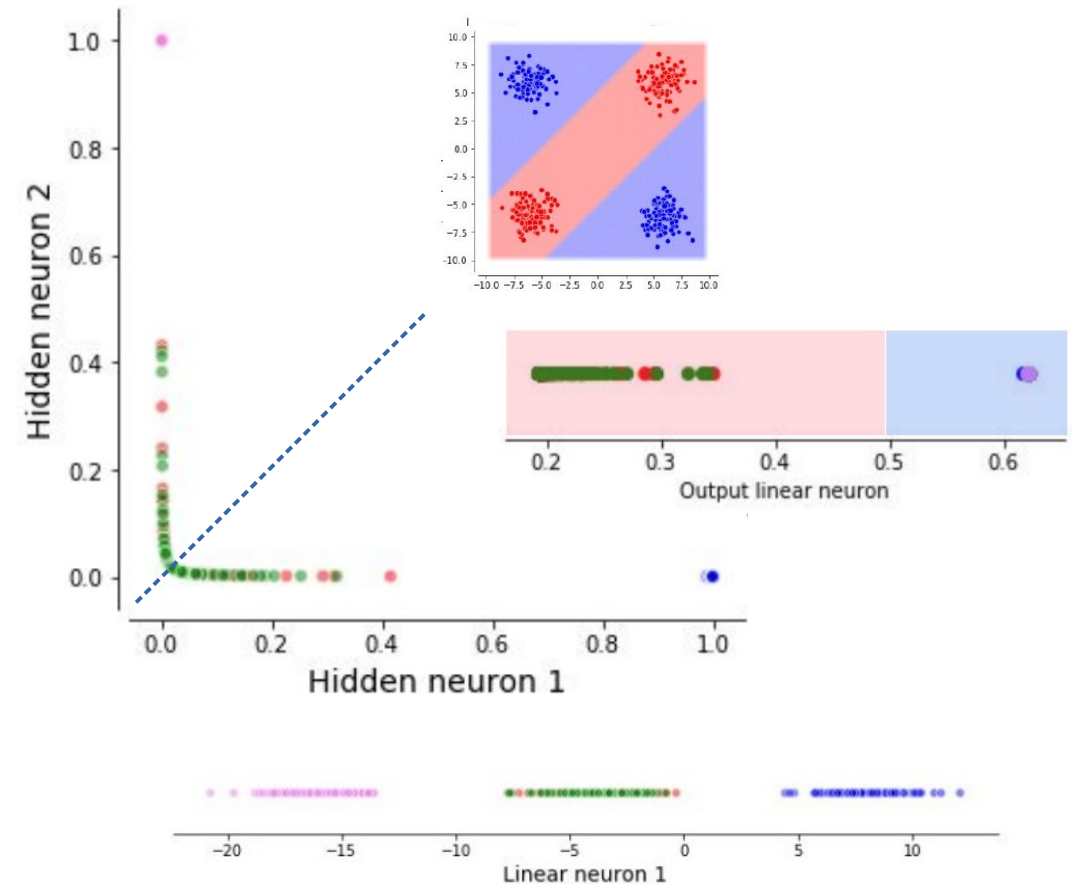


A Two-Layer Neural Network

Going back to the original problem



- Squashing our samples to the range $[0,1]$
- The hidden-layer provides a non-linear input space.



Nice, but ...

What if we have a more complex problem?



A close-up shot from the movie Inception showing Leonardo DiCaprio and Matt Damon. DiCaprio is on the left, looking slightly to the right with a serious expression. Damon is on the right, leaning in towards DiCaprio. The background is blurred, showing what appears to be an office or meeting room setting.

WE NEED TO GO

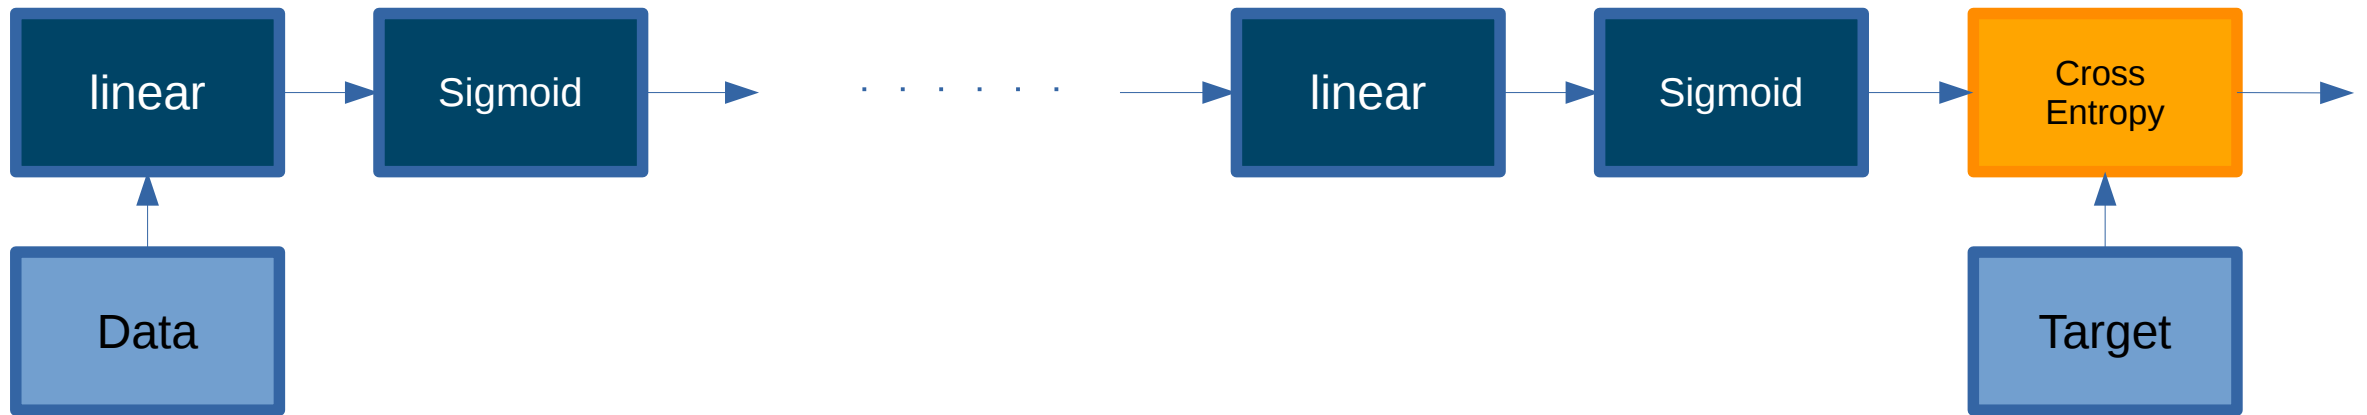
DEEPER

Deep Neural Networks

[adding more and more layers | something something “deep learning”]


Deep Neural Networks

Further extending the previous schematic



Deep Neural Networks


Activation Function – Rectifier Linear Unit



$$f_{relu} = \max(0, x)$$

Deep Neural Networks

Activation Function – Rectifier Linear Unit



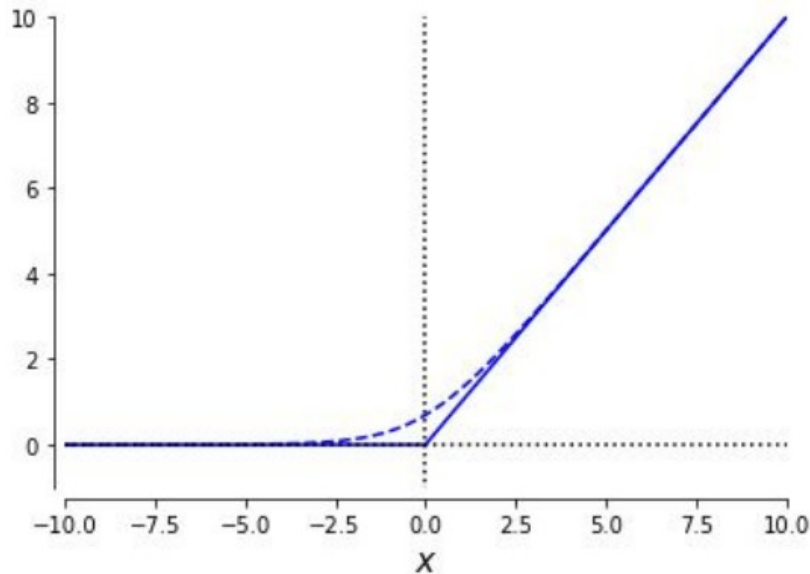
Characteristics:

- Point-wise operation
- Not linear, but piece-wise linear
- Cut the space into polyhedra

$$f_{relu} = \max(0, x)$$

Deep Neural Networks

Activation Function – Rectifier Linear Unit



$$f_{relu} = \max(0, x)$$

Characteristics:

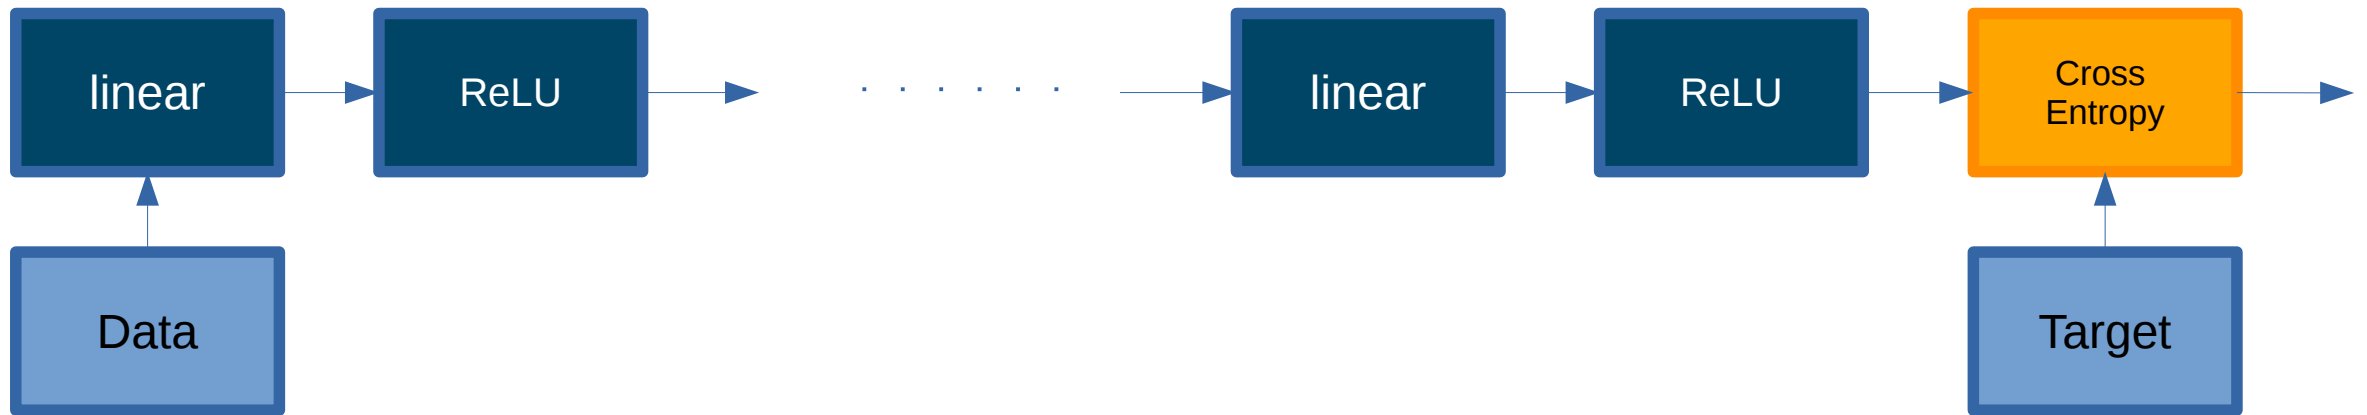
- Point-wise operation
- Not linear, but piece-wise linear
- Cut the space into polyhedra

Note

- Dead neurons can occur
- Not differentiable at 0
- Derivatives do not vanish

Deep Neural Networks

Further extending the previous schematic



Nice, but ...

Does it always work?



Deep Neural Networks

Universal Approximation Theorem [Cybenko, 1989]

Given a continuous function from the hypercube to a single real value.

A large network **can approximate** (up to some error epsilon), **not represent**, any smooth function.

Deep Neural Networks

Universal Approximation Theorem [Cybenko, 1989]

Given a continuous function from the hypercube to a single real value.

A large network **can approximate** (up to some error epsilon), **not represent**, any smooth function.

Does not provides guarantees over the “learnability” of such network.

Size of the network grows exponentially w.r.t. the input dimensions

Deep Neural Networks

Universal Approximation Theorem [Cybenko, 1989]

Given a continuous function from the hypercube to a single real value.

A large network **can approximate** (up to some error epsilon), **not represent**, any smooth function.

Does not provides guarantees over the “learnability” of such network.

Size of the network grows exponentially w.r.t. the input dimensions


[Hornik, 1991]

The key is that the stacked components are non-constant and bounded

Deep Neural Networks

Universal Approximation Theorem [Cybenko, 1989]

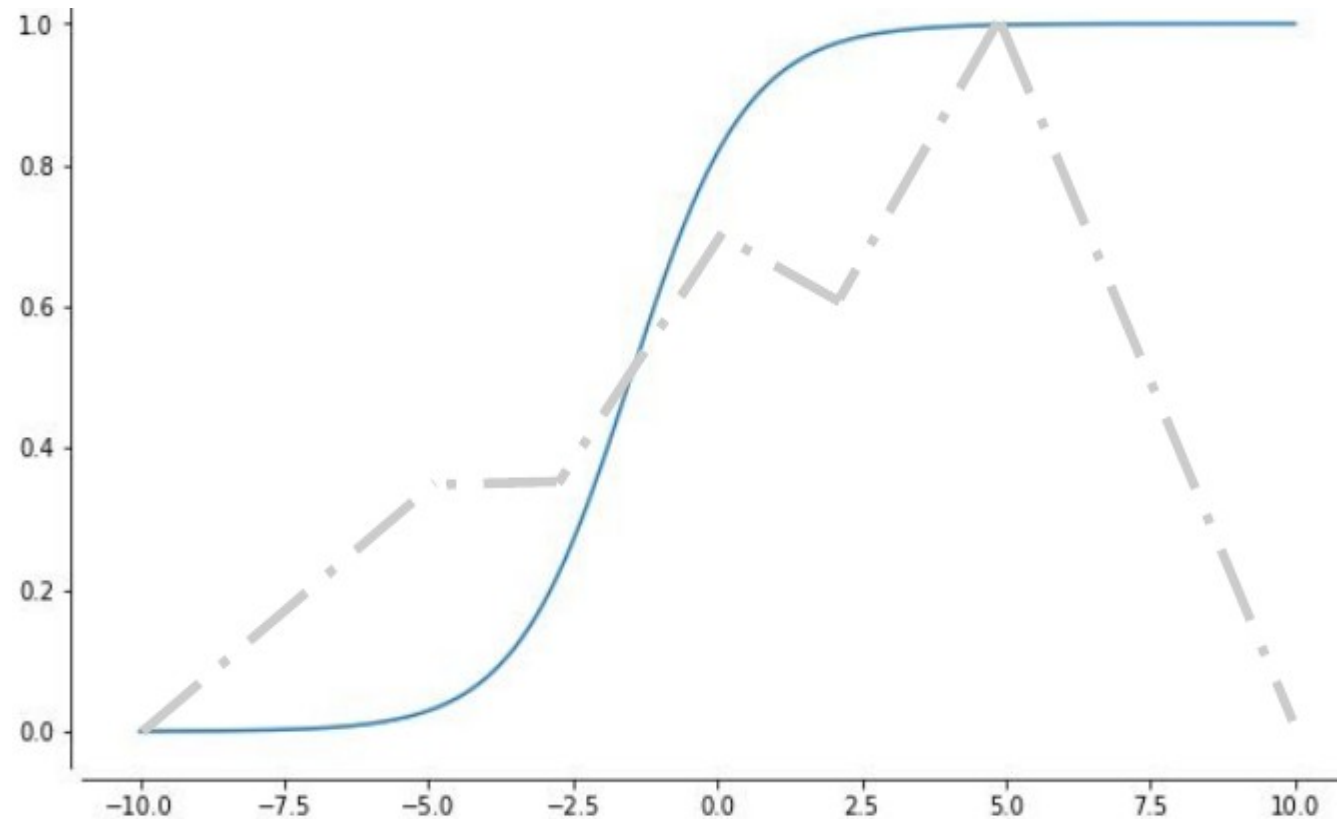
- An intuition on how it works



Deep Neural Networks

Universal Approximation Theorem [Cybenko, 1989]

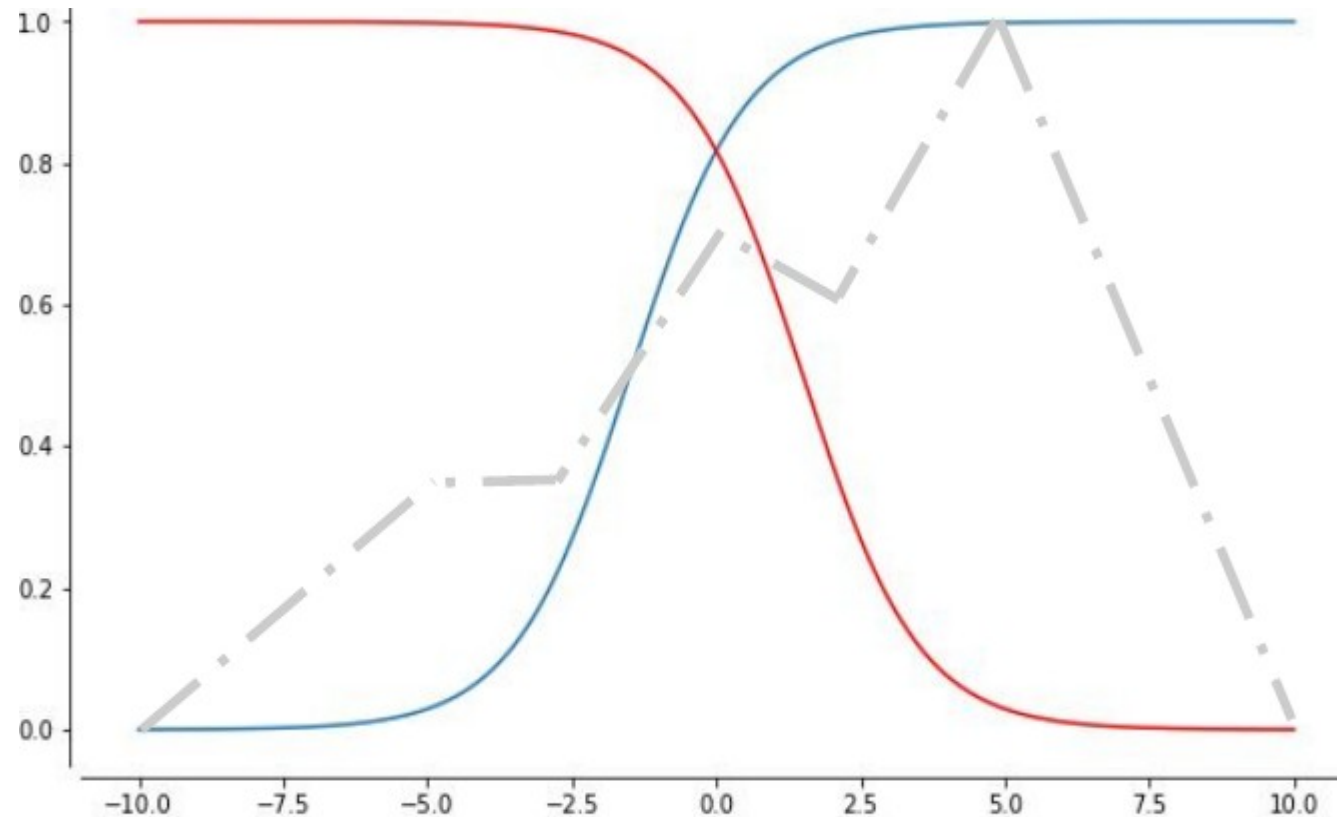
- An intuition on how it works



Deep Neural Networks

Universal Approximation Theorem [Cybenko, 1989]

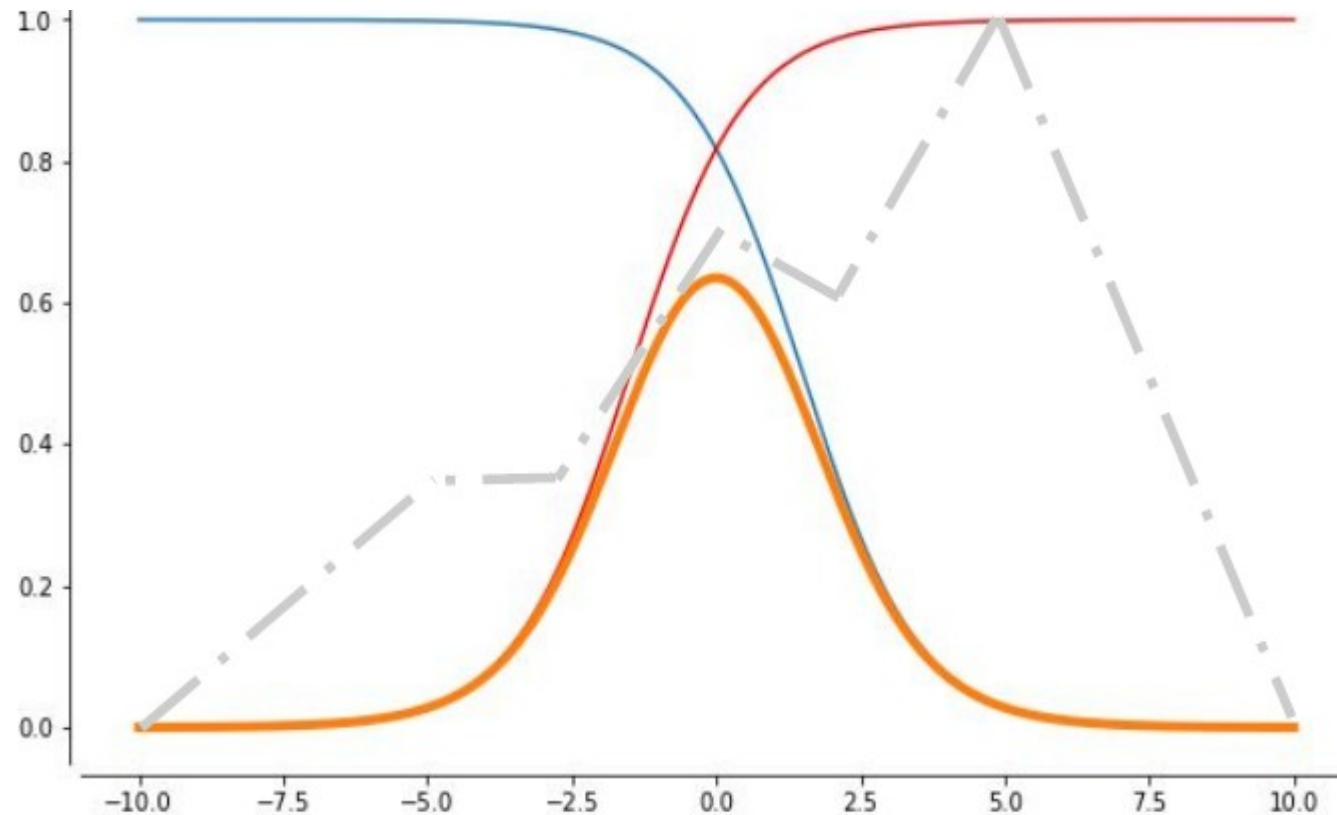
- An intuition on how it works



Deep Neural Networks

Universal Approximation Theorem [Cybenko, 1989]


- An intuition on how it works



Deep Neural Networks

Universal Approximation Theorem [Cybenko, 1989]

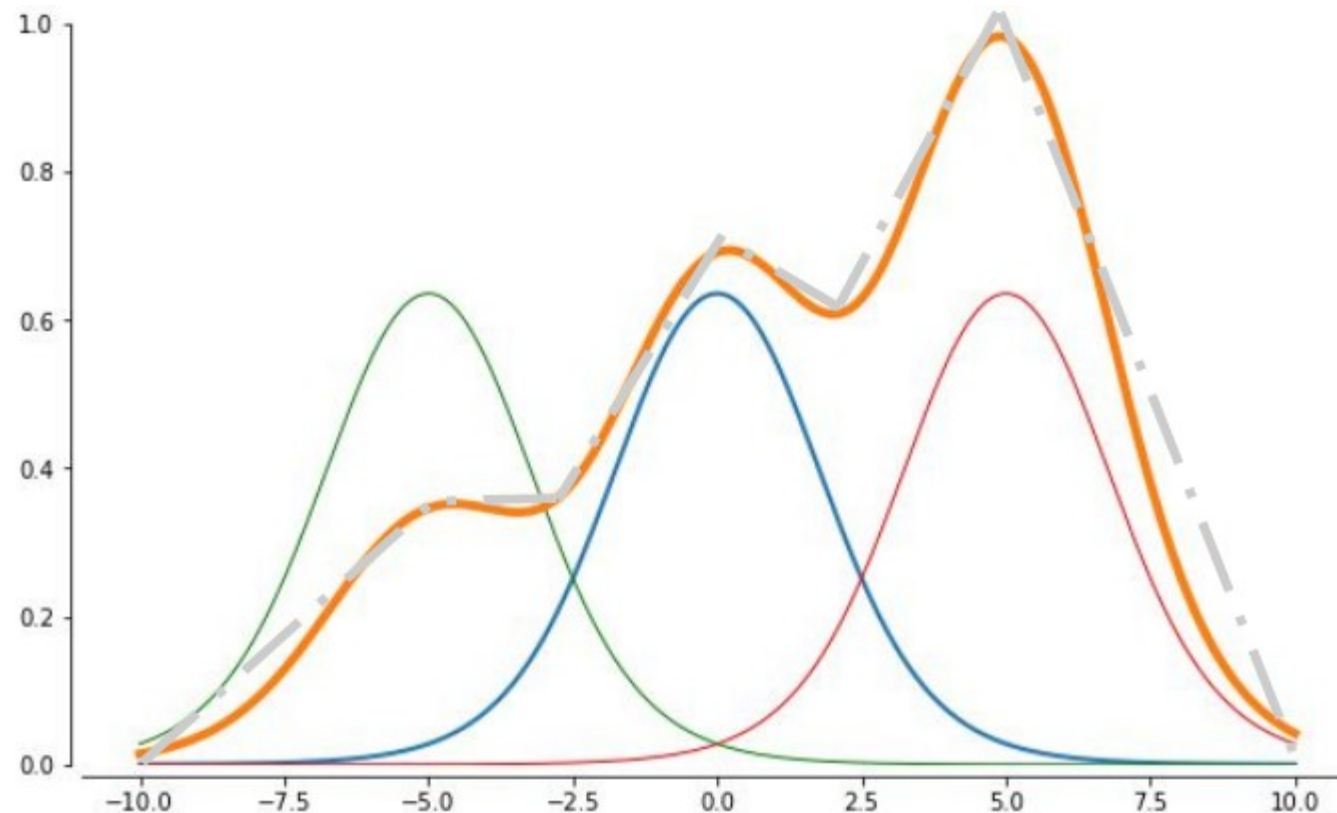
- An intuition on how it works



Deep Neural Networks

Universal Approximation Theorem [Cybenko, 1989]

- An intuition on how it works



Nice, but ...


**What happens in
high-dimensional
spaces?**



Deep Neural Networks

Universal Approximation Theorem [Cybenko, 1989]

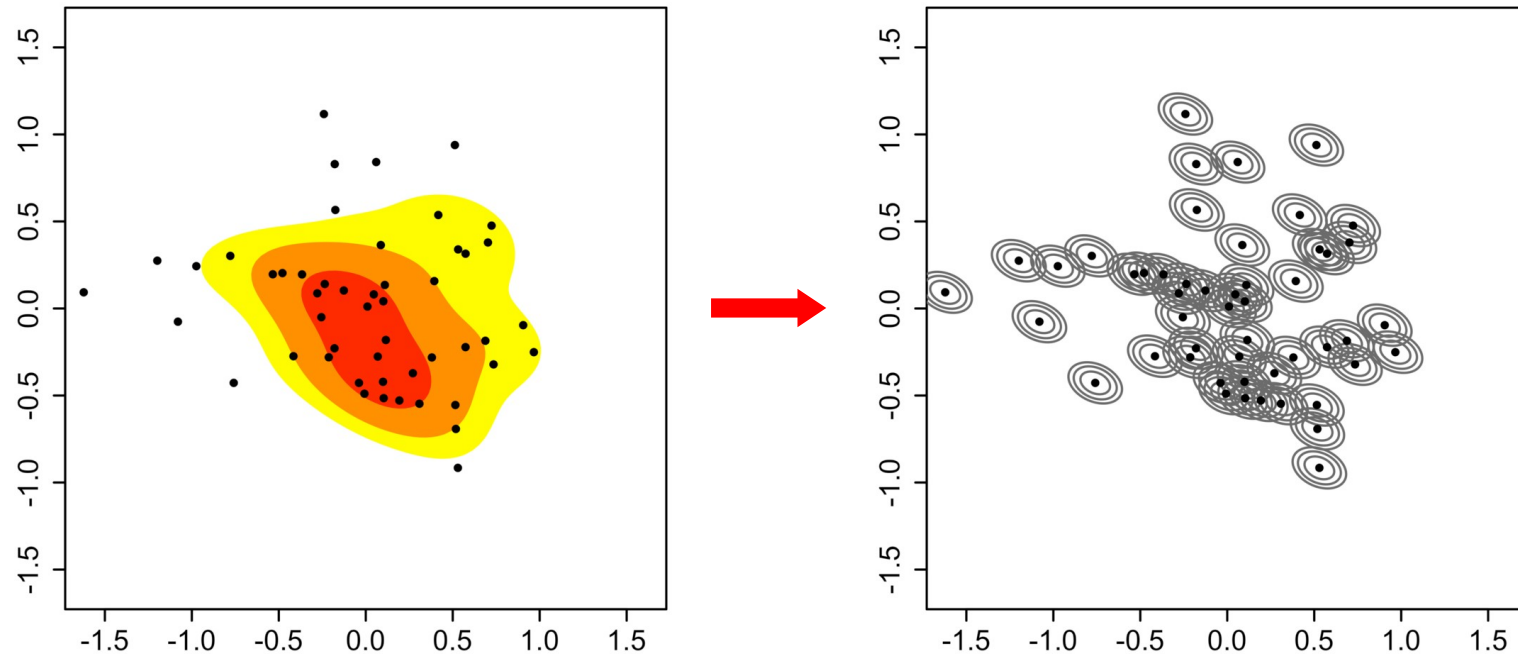
- An intuition on how it works – High-Dimensional Spaces



Deep Neural Networks

Universal Approximation Theorem [Cybenko, 1989]


- An intuition on how it works – High-Dimensional Spaces



Deep Neural Networks

Universal Approximation Theorem [Cybenko, 1989]

- An intuition on how it works – High-Dimensional Spaces



Ok, but ...

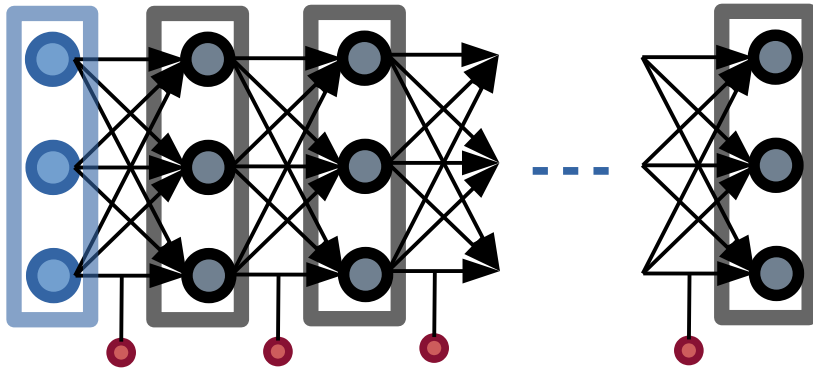
**Why deeper rather
than wider?**



Deep Neural Networks


Deeper VS. Wider Architectures

Deeper



V
S

Wider



Growth of the partitioning space

- Exponential by depth
- Polynomial by width

Deep Neural Networks

Deeper VS. Wider Architectures



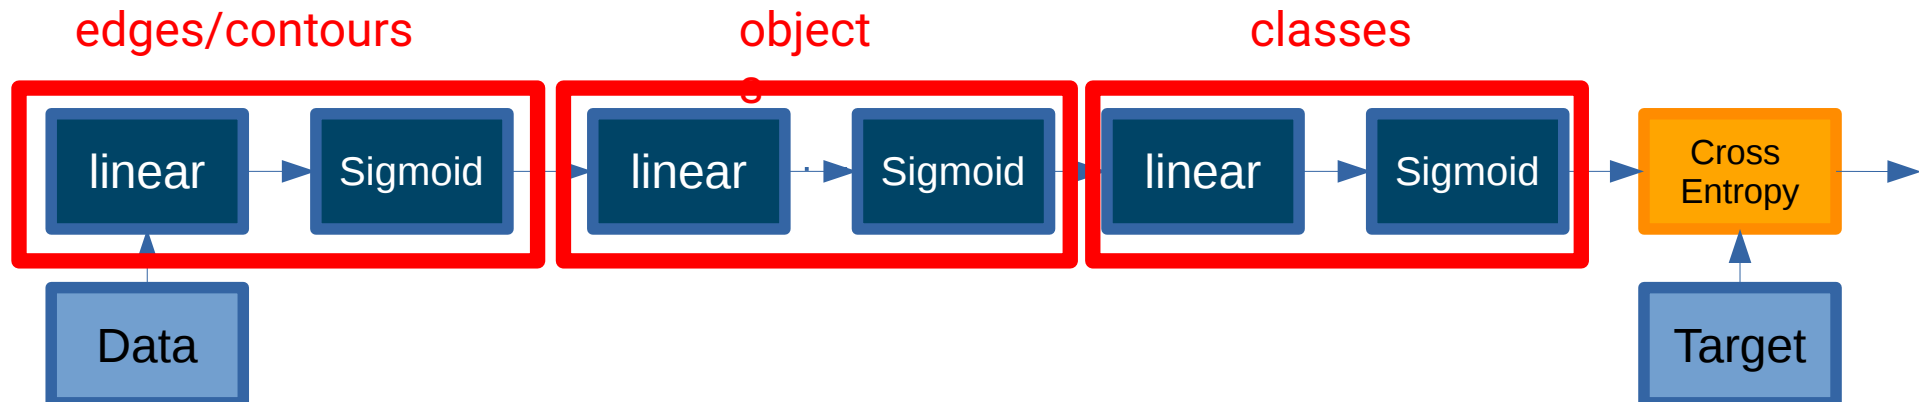
← Learning how to recognize a bicycle

Deep Neural Networks

Deeper VS. Wider Architectures



← Learning how to recognize a bicycle




Learning

[with few layers]


Learning

Process Overview



Learning

Gradient Descent



Learning

Gradient Descent – Algebraic Foundations

$$y = f(\boldsymbol{x}) : \mathbb{R}^d \rightarrow \mathbb{R}$$

$$\frac{\partial y}{\partial \boldsymbol{x}} = \nabla_{\boldsymbol{x}} f(\boldsymbol{x}) = \left[\frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_d} \right]$$

Gradient

Learning

Gradient Descent – Algebraic Foundations

$$y = f(\mathbf{x}) : \mathbb{R}^d \rightarrow \mathbb{R}$$

$$\frac{\partial y}{\partial \mathbf{x}} = \nabla_{\mathbf{x}} f(\mathbf{x}) = \left[\frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_d} \right]$$

Gradient

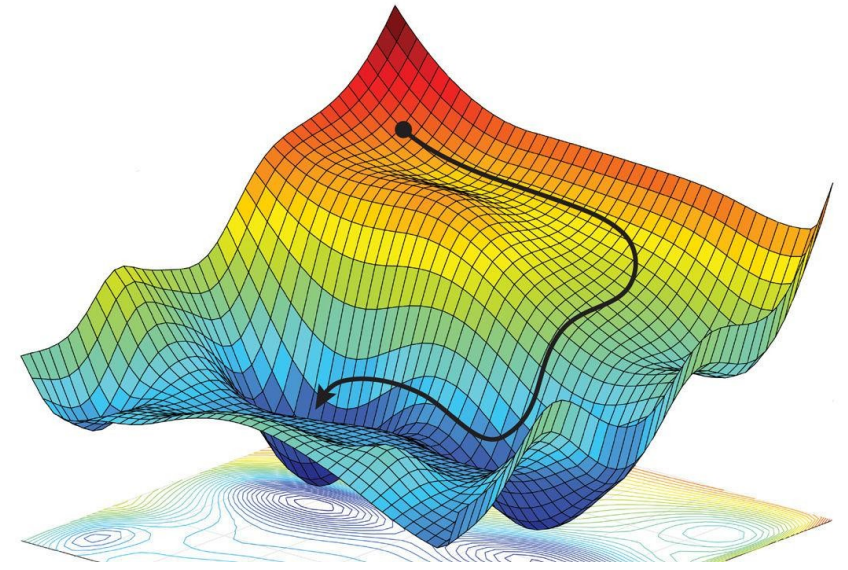
$$\frac{\partial y}{\partial \mathbf{x}} = \mathbf{J}_{\mathbf{x}} f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_d} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_k}{\partial x_1} & \cdots & \frac{\partial f_k}{\partial x_d} \end{bmatrix}$$

Jacobian

Learning

Gradient Descent

$$\theta_{t+1} = \theta_t - \alpha_t \nabla_{\theta} L(\theta_t)$$



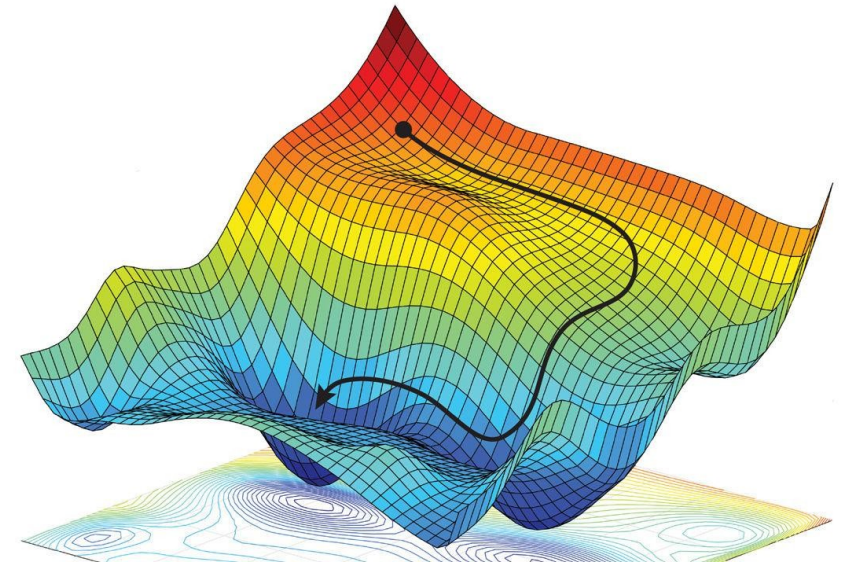
[doi:10.1126/science.aau0577]

Learning

Gradient Descent

$$\theta_{t+1} = \theta_t - \alpha_t \nabla_{\theta} L(\theta_t)$$

$$\begin{aligned} \nabla_{\theta} L(\theta_t) &= \nabla_{\theta} \sum_i l(f(x^{(i)}, \theta_t), y^{(i)}) \\ &= \sum_i \nabla_{\theta} l(f(x^{(i)}, \theta_t), y^{(i)}) \end{aligned}$$



[doi:10.1126/science.aau0577]

Learning


Gradient Descent

$$\theta_{t+1} = \theta_t - \alpha_t \nabla_{\theta} L(\theta_t)$$

$$\begin{aligned} \nabla_{\theta} L(\theta_t) &= \nabla_{\theta} \sum_i l(f(x^{(i)}, \theta_t), y^{(i)}) \\ &= \sum_i \nabla_{\theta} l(f(x^{(i)}, \theta_t), y^{(i)}) \end{aligned}$$

Characteristics:

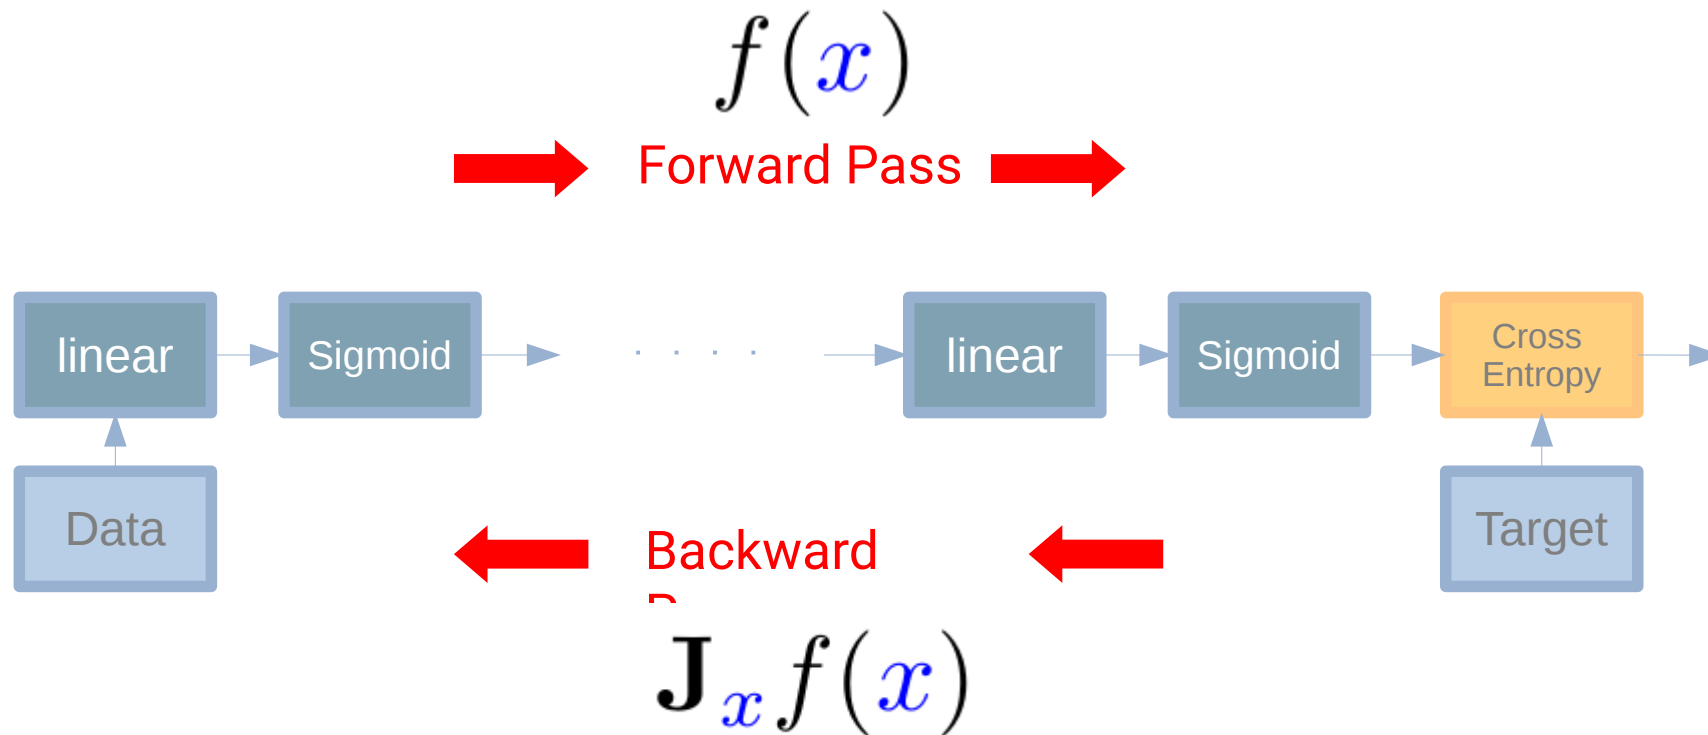
- Works for any smooth function
- Less guarantees for some non-smooth targets
- Converges to local optimum
- Critical effect of the *Learning rate*



[doi:10.1126/science.aau0577]


Learning

Process Overview



Learning

Back-Propagation Algorithm

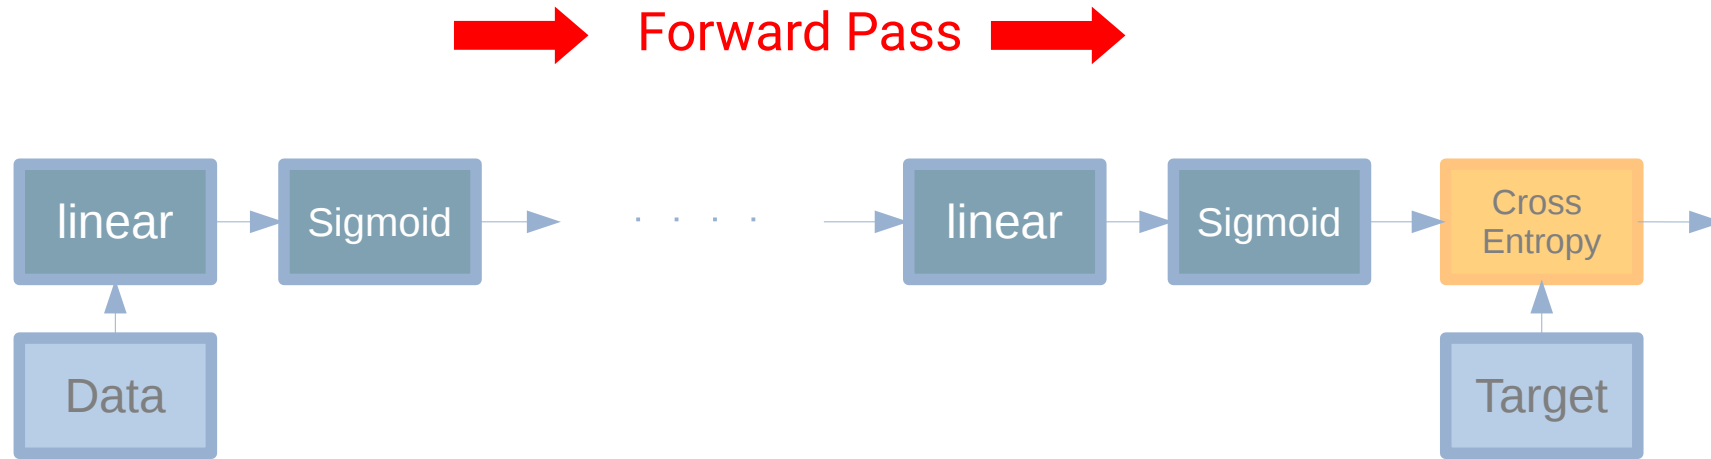


$$f_1(x)$$

Looking the forward pass as a composition

Learning

Back-Propagation Algorithm




$$f_2(f_1(x))$$

Looking the forward pass as a composition

Learning

Back-Propagation Algorithm

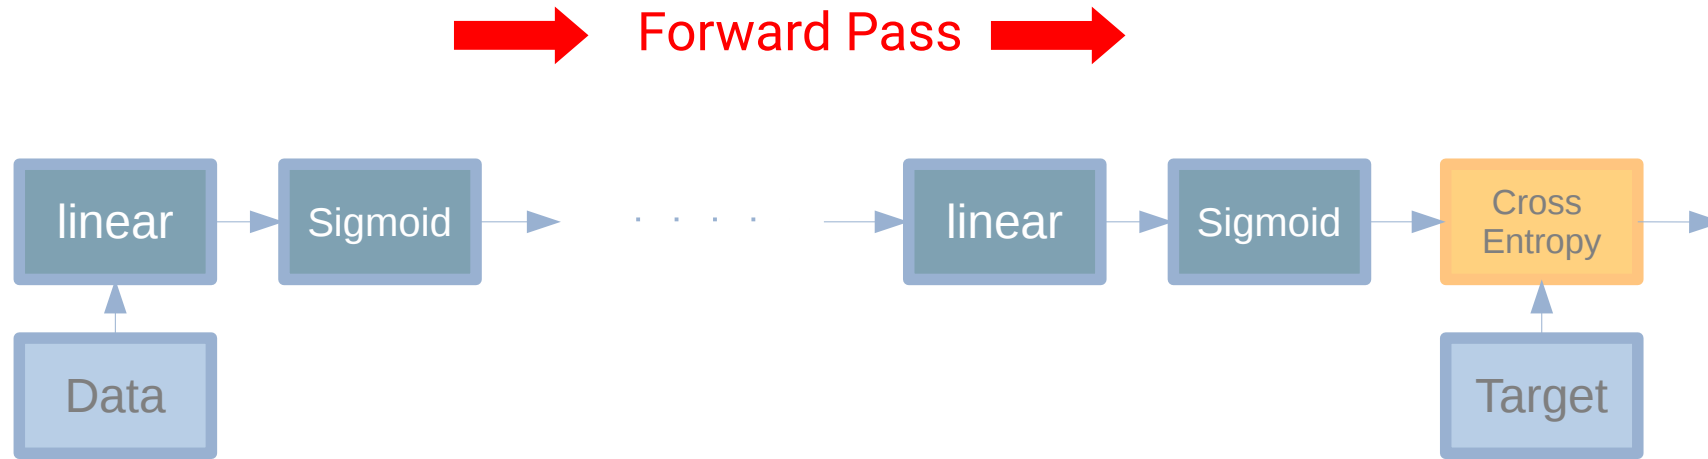


$$f_{k-1} \left(f_{k-2} \left(\dots f_2 \left(f_1(x) \right) \right) \right)$$

Looking the forward pass as a composition

Learning

Back-Propagation Algorithm




$$y = f_k(f_{k-1}(f_{k-2}(\dots f_2(f_1(\textcolor{blue}{x}))))))$$

Looking the forward pass as a composition


Learning

Back-Propagation Algorithm



Learning


Back-Propagation Algorithm



$$y = f(g(x)) \frac{\partial y}{\partial x}$$

Learning

Back-Propagation Algorithm

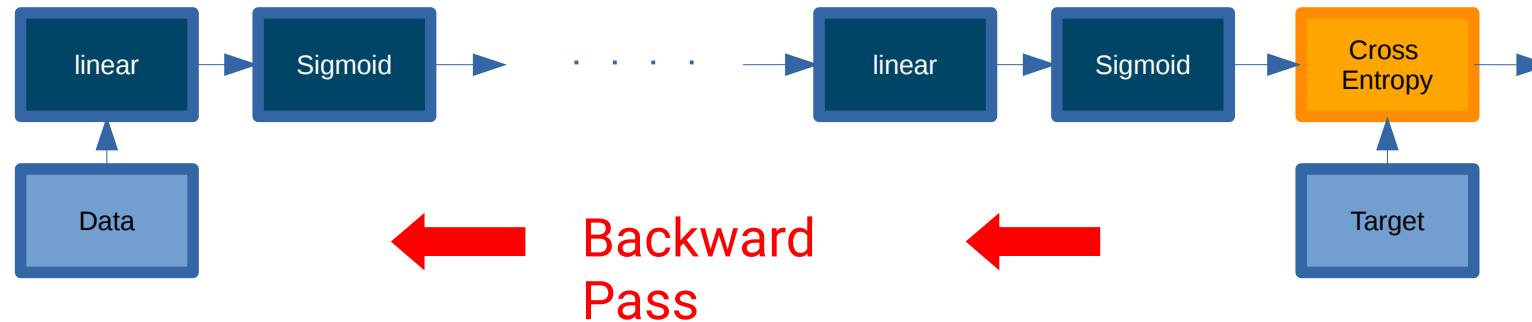


$$y = f(g(x)) \frac{\partial y}{\partial x} = \frac{\partial y}{\partial g} \frac{\partial g}{\partial x}$$

$$\frac{dL}{dw} = \frac{dL}{dy} \cdot \frac{dy}{dz} \cdot \frac{dz}{dw}$$

Learning

Back-Propagation Algorithm



$$y = f(g(x)) \frac{\partial y}{\partial x} = \frac{\partial y}{\partial g} \frac{\partial g}{\partial x}$$

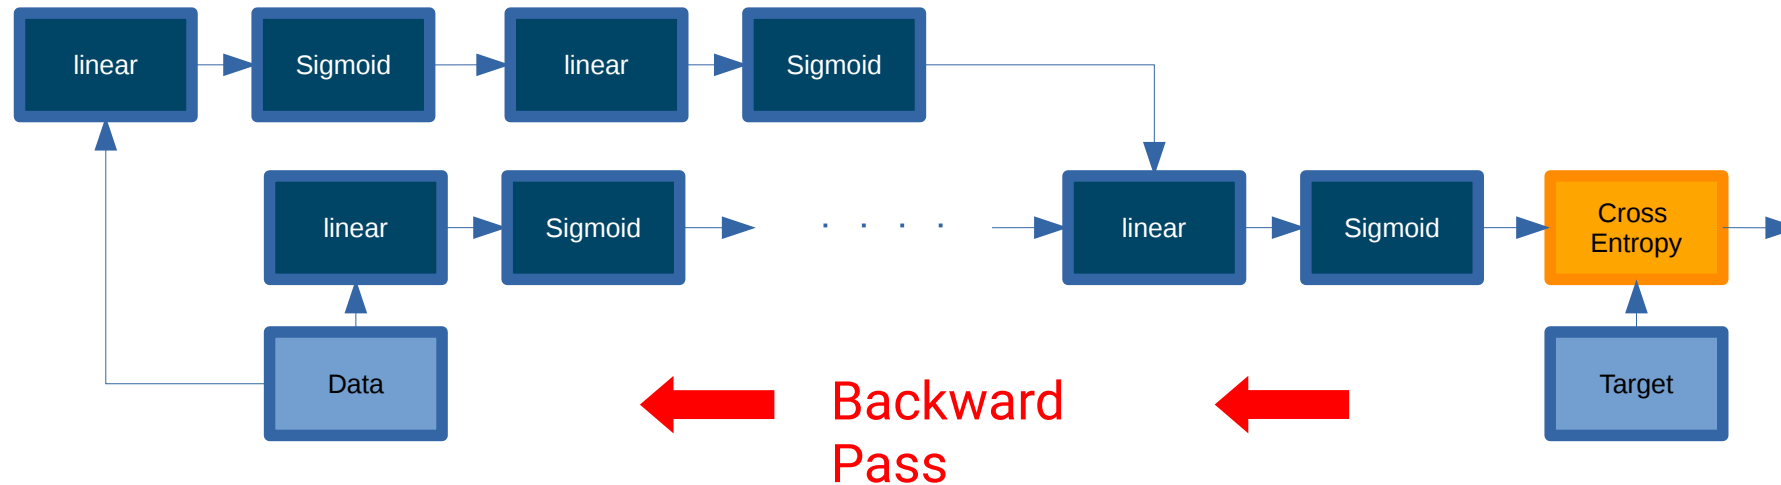
$$\frac{dL}{dw} = \frac{dL}{dy} \cdot \frac{dy}{dz} \cdot \frac{dz}{dw}$$

Characteristics:

- Chain rule for derivative computation
- Computations can be re-used (makes is linear [faster], otherwise quadratic)

Learning

Back-Propagation Algorithm



$$y = f(g(x)) \frac{\partial y}{\partial x} = \frac{\partial y}{\partial g} \frac{\partial g}{\partial x}$$

$$y = f(g(X)) \frac{\partial y}{\partial X} = \sum_{i=1}^m \frac{\partial y}{\partial g^{(i)}} \frac{\partial g^{(i)}}{\partial X}$$

Characteristics:

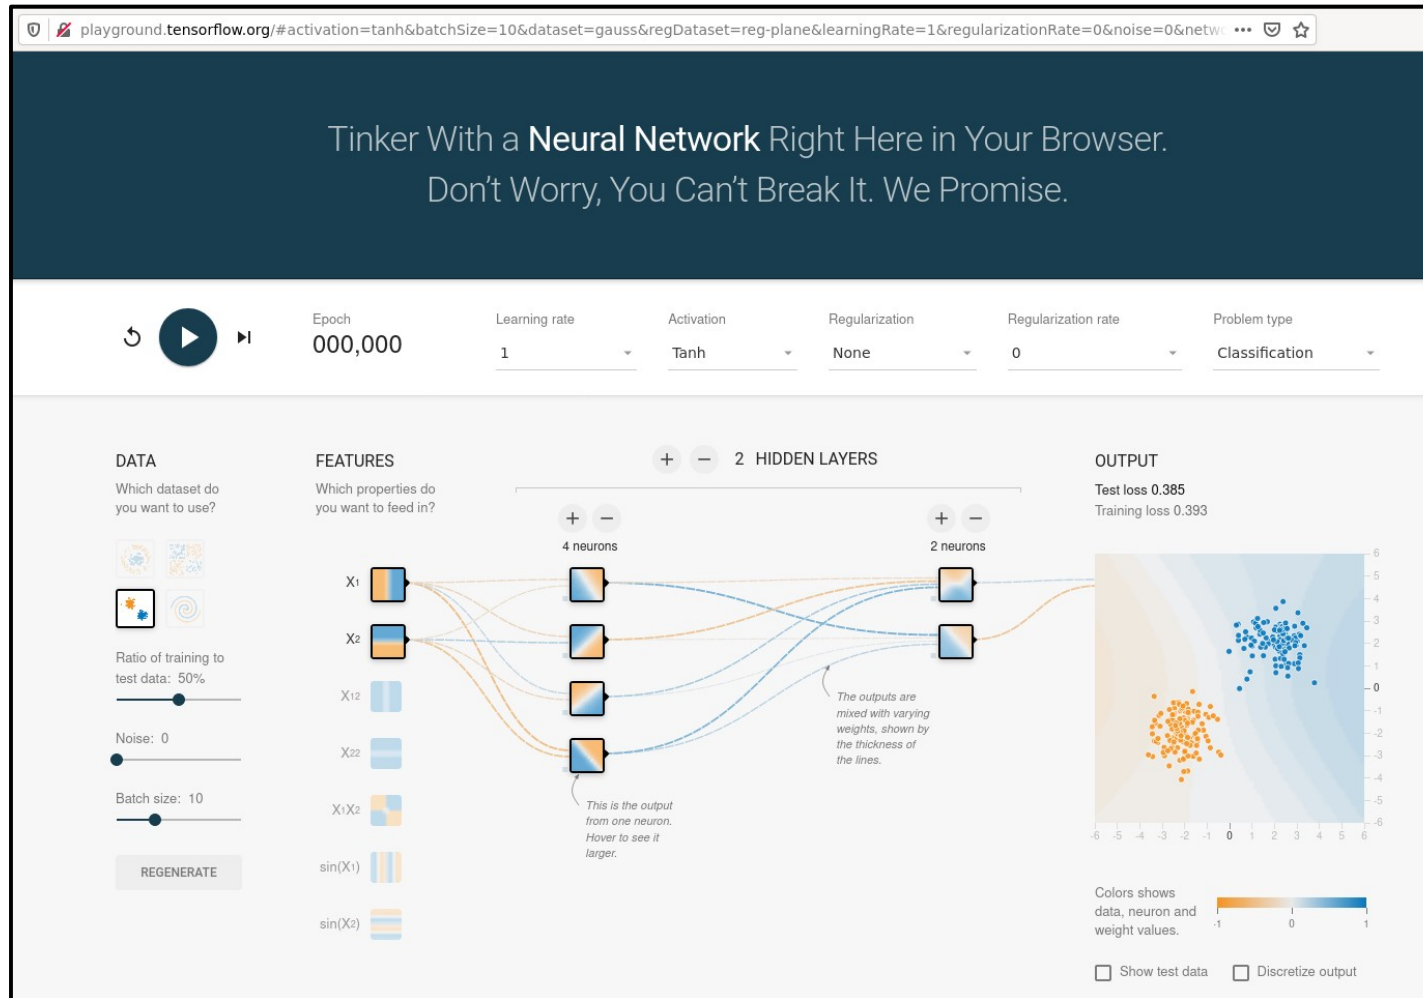
- Chain rule for derivative computation
- Computations can be re-used (makes is linear [faster], otherwise quadratic)

wait, ...
That's it?



From Shallow to Deep Neural Networks

Some Extra Practice – <http://playground.tensorflow.org>



Summarizing

[Finally :D]

Summarizing

- **A From Neurons to Networks**
 - Neurons → Layers → Networks

Summarizing

- **A From Neurons to Networks**

- Neurons → Layers → Networks

- **Power through Composition**

- It is not about a single unit but their combination
- Capable of approximating any function producing a single real value as output
- Better deeper (exponential) than wider (polynomial) architectures

Summarizing

- **A From Neurons to Networks**

- Neurons → Layers → Networks

- **Power through Composition**

- It is not about a single unit but their combination
- Capable of approximating any function producing a single real value as output
- Better deeper (exponential) than wider (polynomial) architectures

- **Some Enablers**

- Efficient algorithmic computations
- Use of dedicated hardware

Pay Attention...

[one last tip for today]

Pay attention to...

- Everything



References

Universal Approximation Capabilities of Deep Neural Networks

- G. Cybenko, Approximation by superpositions of a sigmoidal function, Math. Control Signals Systems, 2 (1989), 303–314.
<https://link.springer.com/article/10.1007/BF02134016>
- K. Hornik Approximation Capabilities of Multilayer Feedforward Networks
https://web.njit.edu/~usman/courses/cs675_spring20/hornik-nn-1991.pdf

Deep VS. Wide architectures

- Guido Montúfar, Razvan Pascanu, Kyunghyun Cho, Yoshua Bengio, On the Number of Linear Regions of Deep Neural Networks. NeurIPS 2014
<https://papers.nips.cc/paper/2014/file/109d2dd3608f669ca17920c511c2a41e-Paper.pdf>

ReLU

- R H Hahnloser¹, R Sarpeshkar, M A Mahowald, R J Douglas, H S Seung. Digital selection and analogue amplification coexist in a cortex-inspired silicon circuit. Nature 2000
<https://pubmed.ncbi.nlm.nih.gov/10879535/>

Questions?



From Shallow to Deep Neural Networks

[Building More Complex Models]

José Oramas